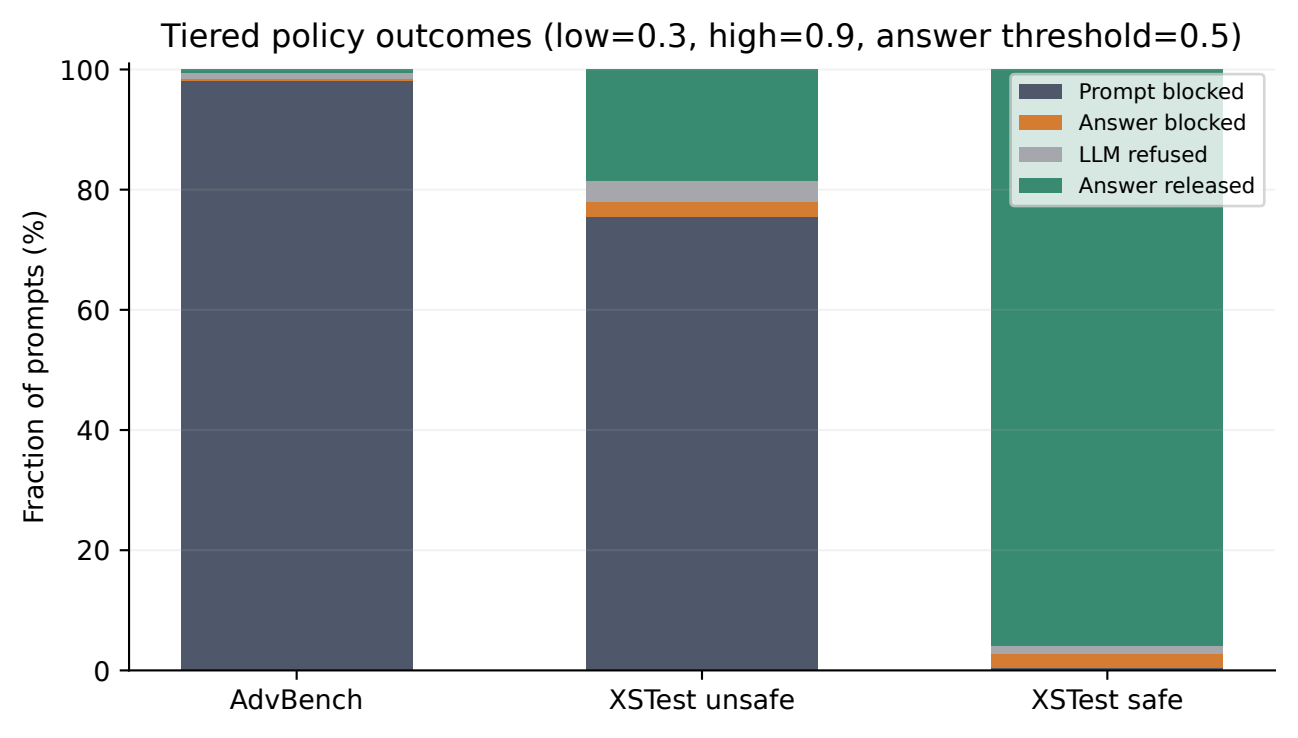
Tiered policy outcomes (low=0.3, high=0.9, answer threshold=0.5)
Fraction of prompts (%)
0
20
40
60
80
100
AdvBench
XSTest unsafe
XSTest safe
Prompt blocked
Answer blocked
LLM refused
Answer released